

Introduction

- **Scope:** Detailed study of CUDA-based matrix multiplication performance across six kernel versions.
- **Metrics:** Execution time and GFLOPS across varying matrix sizes.
- **Objective:** Analyze how each optimization scales and impacts performance.

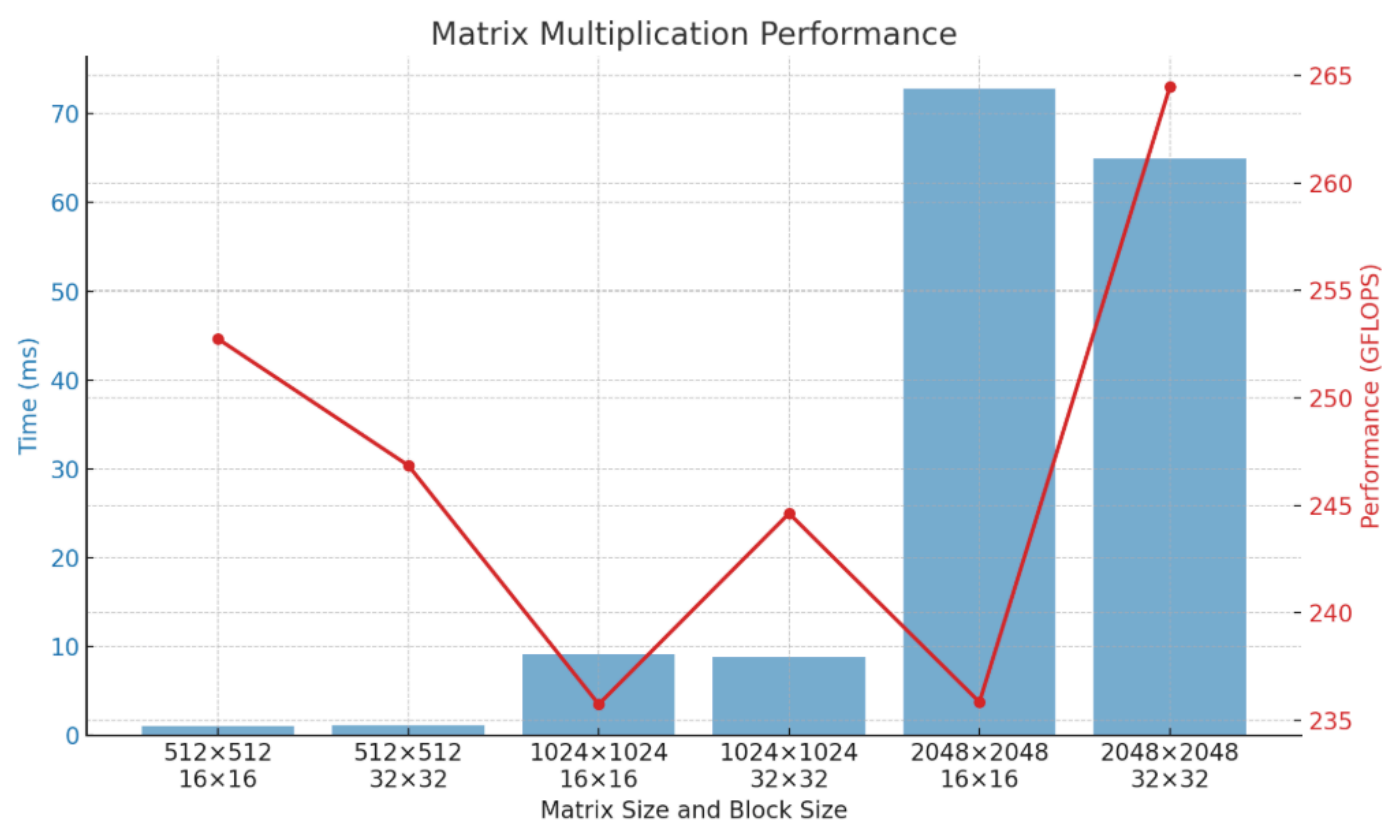
Device Specifications

Property	Value
Total Global Memory	8 106 MB
Shared Memory per Block	48 KB
Registers per Block	65 536
Warp Size	32
Max Threads per Block	1 024
Multi-Processor Count (SMs)	14
Max threads per multiprocessor	: 2048

V1 – Baseline

Result

Performance table and plot to be inserted.



- The blue bars show the Time (ms) for each matrix size and block size.
- The red line shows the corresponding GFLOPS performance.

Matrix Size	Block Size	Time (ms)	GFLOPS
512x512	16x16	1.062	252.75
512x512	32x32	1.087	246.86
1024x1024	16x16	9.109	235.75
1024x1024	32x32	8.779	244.63
2048x2048	16x16	72.842	235.85
2048x2048	32x32	64.962	264.46

Code Snippet

```

__global__ void V1_baselineKernel(const float* A, const float* B, float* C, int N) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    if (row < N && col < N) {
        float sum = 0.0f;
        for (int k = 0; k < N; ++k) {
            sum += A[row * N + k] * B[k * N + col];
        }
        C[row * N + col] = sum;
    }
}

```

Technique: Naïve Global Memory Access

- Reads elements directly from global memory for each multiply-add.
- No tiling or caching; simple row-by-column dot-product.

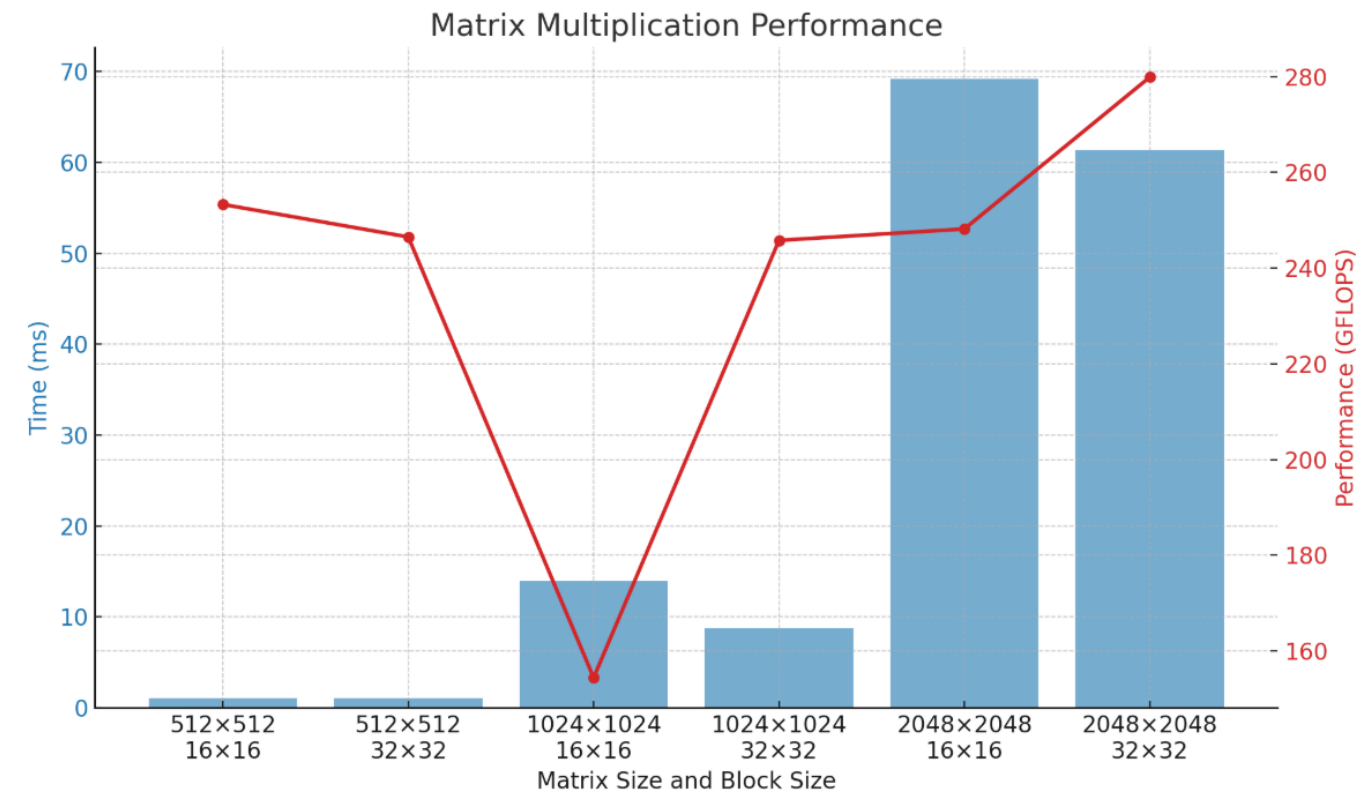
Explanation:

- **Memory-bound:** Frequent reads from global memory dominate latency.
- **Compute underutilization:** Many warps idle waiting for memory.
- **Scales poorly:** Time grows quadratically with N.
- **Baseline for comparison:** Establishes reference GFLOPS (~235–265). Only minimal code complexity.

V2 – Loop Unrolling

Result

Performance table and plot to be inserted.



- The blue bars show the Time (ms) for each matrix size and block size.
- The red line shows the corresponding GFLOPS performance.

Matrix Size	Block Size	Time (ms)	GFLOPS
512×512	16×16	1.060	253.28
512×512	32×32	1.089	246.49
1024×1024	16×16	13.905	154.44
1024×1024	32×32	8.738	245.78
2048×2048	16×16	69.227	248.17
2048×2048	32×32	61.388	279.86

Code Snippet

```
__global__ void V2_loopUnrollKernel(const float* A, const float* B, float* C, int N) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    if (row < N && col < N) {
        float sum = 0.0f;
        int k = 0;
        for (; k <= N - 4; k += 4) {
            sum += A[row * N + k] * B[k * N + col];
            sum += A[row * N + k + 1] * B[(k + 1) * N + col];

```

```

        sum += A[row * N + k + 2] * B[(k + 2) * N + col];
        sum += A[row * N + k + 3] * B[(k + 3) * N + col];
    }
    for (; k < N; ++k) {
        sum += A[row * N + k] * B[k * N + col];
    }
    C[row * N + col] = sum;
}
}

```

Technique: Loop Unrolling

- Unrolls inner `k`-loop by factor of 4 to reduce branch overhead.
- More arithmetic per loop iteration.


Explanation:

- **Reduced loop control overhead:** Fewer branch checks improve throughput.
- **Mixed results:** Small matrices (~512) see negligible change; large (1024,2048) improve slightly when block size=32.
- **Compute vs. memory:** Unrolling helps only when arithmetic latency hides memory fetches; limited by global memory.
- **Inconsistent scaling:** Performance benefit depends on block configuration.

V3 – Shared Memory Tiling

Result

Performance table and plot to be inserted.



- The blue bars show the Time (ms) for each matrix size and block size.
- The red line shows the corresponding GFLOPS performance.

Matrix Size	Block Size	Time (ms)	GFLOPS
512x512	16x16	0.528	507.94
512x512	32x32	0.507	529.57
1024x1024	16x16	4.190	512.58
1024x1024	32x32	3.912	548.89
2048x2048	16x16	33.285	516.15
2048x2048	32x32	42.425	404.94

Code Snippet

```
template <int TILE_SIZE>
__global__ void V3_sharedMemoryKernel(const float* A, const float* B, float* C, int N) {
    __shared__ float As[TILE_SIZE][TILE_SIZE];
    __shared__ float Bs[TILE_SIZE][TILE_SIZE];
    int row = blockIdx.y * TILE_SIZE + threadIdx.y;
    int col = blockIdx.x * TILE_SIZE + threadIdx.x;
    float sum = 0.0f;
    for (int t = 0; t < (N + TILE_SIZE - 1) / TILE_SIZE; ++t) {
        // Load tiles
        As[threadIdx.y][threadIdx.x] = (row < N && t*TILE_SIZE+threadIdx.x < N)
```

```

        ? A[row*N + t*TILE_SIZE + threadIdx.x] : 0.0f;
Bs[threadIdx.y][threadIdx.x] = (col < N && t*TILE_SIZE+threadIdx.y < N)
        ? B[(t*TILE_SIZE + threadIdx.y)*N + col] : 0.0f;
__syncthreads();
for (int k=0; k < TILE_SIZE; ++k) {
    sum += As[threadIdx.y][k] * Bs[k][threadIdx.x];
}
__syncthreads();
}
if (row < N && col < N) C[row*N + col] = sum;
}

```

Technique: Shared Memory Tiling

- Loads sub-blocks (tiles) of A and B into fast shared memory.
- Reuses each tile across TILE_SIZE iterations.
- Synchronizes with `__syncthreads()` to ensure complete tile loads.

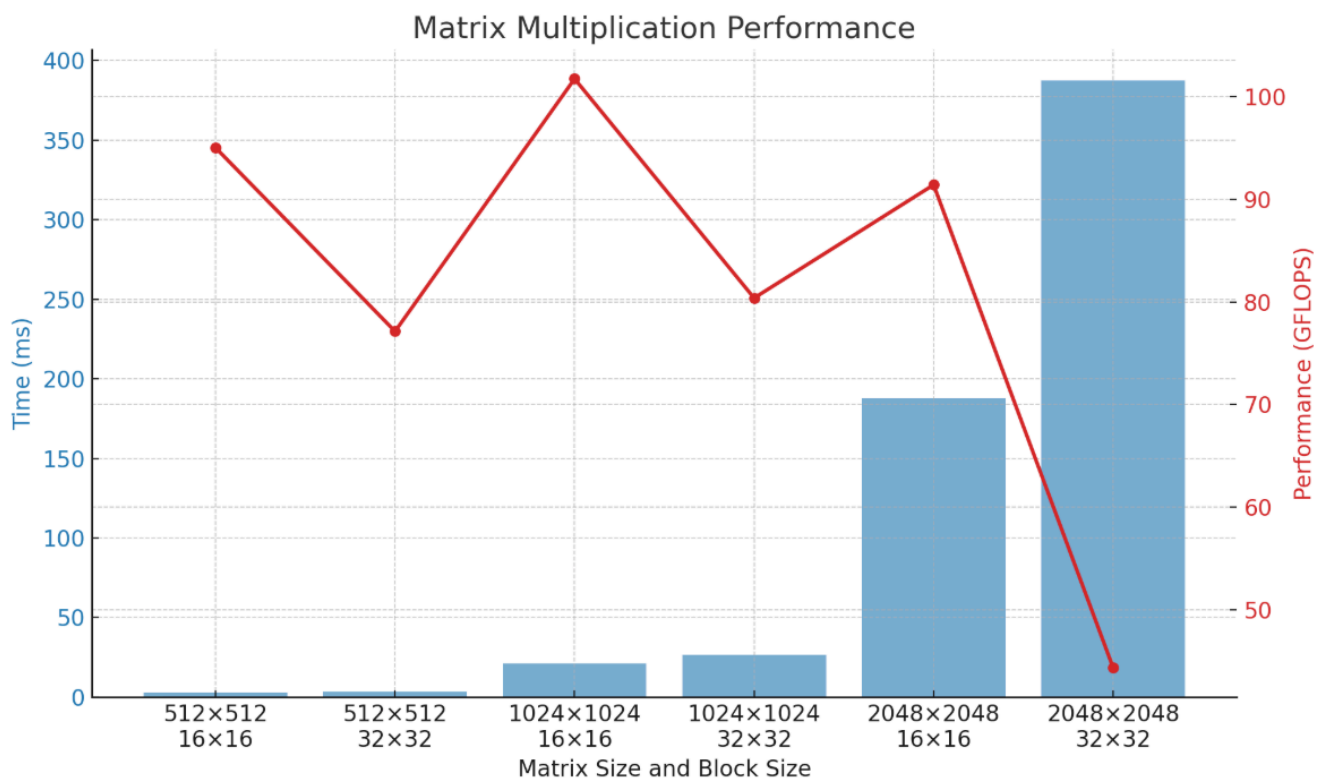
Explanation:

- **Memory coalescing:** Bulk reads from global memory amortized over many arithmetic operations.
- **Latency hiding:** Shared memory (~100× faster) reduces global loads.
- **High throughput:** Achieves >500 GFLOPS on 512–1024 sizes.
- **Block size sensitivity:** 32×32 best for mid-sizes; large matrices see diminishing shared-memory reuse or increased synchronization cost.

V4 – Thread Coarsening (Coarse-Grained)

Result

Performance table and plot to be inserted.



- The blue bars show the Time (ms) for each matrix size and block size.
- The red line shows the corresponding GFLOPS performance.

Matrix Size	Block Size	Time (ms)	GFLOPS
512x512	16x16	2.824	95.04
512x512	32x32	3.480	77.13
1024x1024	16x16	21.103	101.76
1024x1024	32x32	26.715	80.38
2048x2048	16x16	187.926	91.42
2048x2048	32x32	387.539	44.33

Code Snippet

```
__global__ void V4_threadCoarseningKernel(const float* A, const float* B, float* C, int N) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col_start = (blockIdx.x * blockDim.x + threadIdx.x) * COARSE_FACTOR;
    if (row < N) {
        for (int c=0; c<COARSE_FACTOR; ++c) {
            int col = col_start + c;
            if (col < N) {
                float sum = 0.0f;
                for (int k=0; k<N; ++k)
```



```
        sum += A[row*N + k] * B[k*N + col];
        C[row*N + col] = sum;
    }
}
```

Technique: Thread Coarsening

- Each thread computes multiple output elements (`COARSE_FACTOR`).
- Reduces launch overhead and increases per-thread workload.

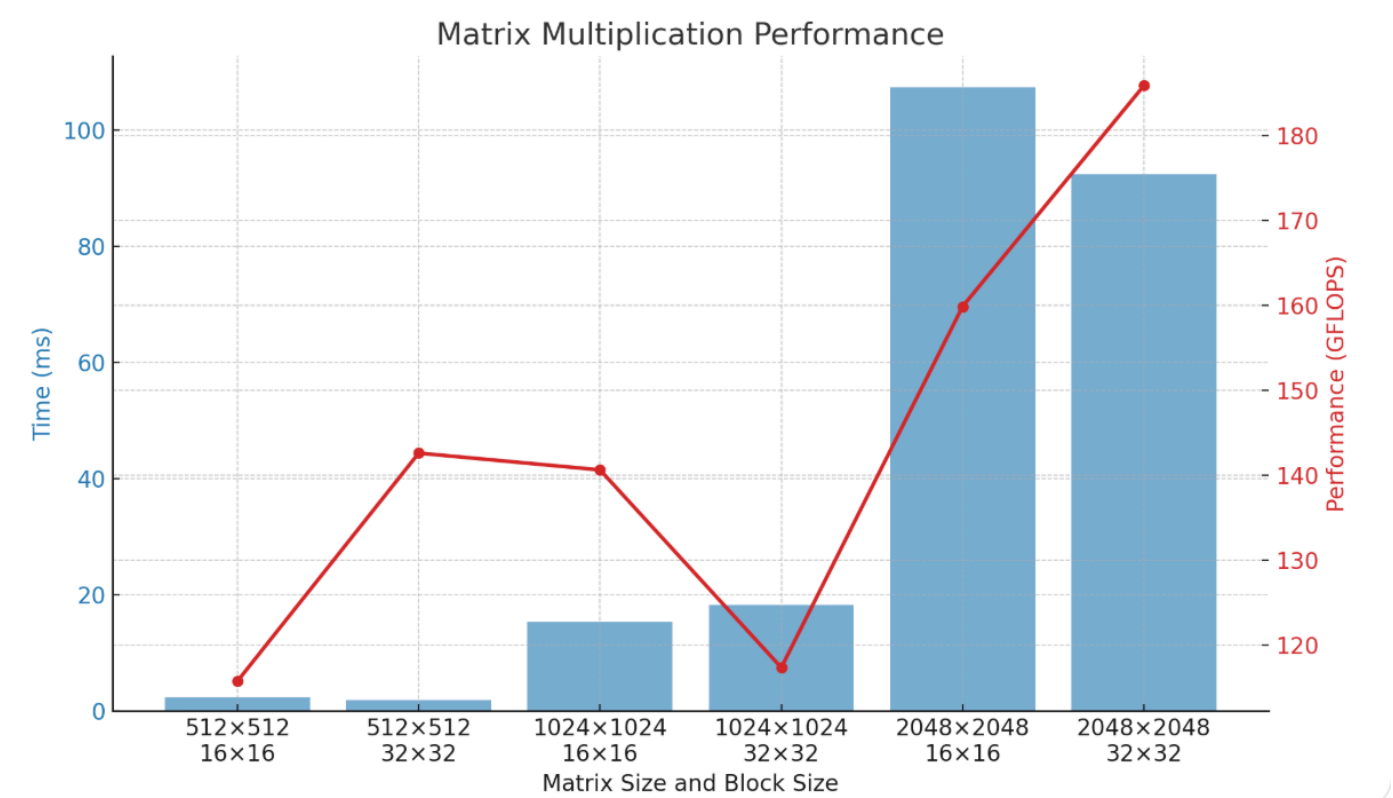
Explanation:

- **Increased register pressure:** More partial sums in registers leading to spills.
- **Poor memory locality:** Each thread loads disparate B elements, harming coalescing.
- **Underutilization:** Many threads idle on heavy FMAs; global loads dominate.
- **Overall slowdown:** Performance falls below baseline

V5 – Privatization (Register Tiling)

Result

Performance table and plot to be inserted.



- The blue bars show the Time (ms) for each matrix size and block size.

- The red line shows the corresponding GFLOPS performance.

Reg_tile_size = 4

Matrix Size	Block Size	Time (ms)	GFLOPS
512×512	16×16	2.319	115.75
512×512	32×32	1.882	142.60
1024×1024	16×16	15.272	140.62
1024×1024	32×32	18.299	117.35
2048×2048	16×16	107.467	159.86
2048×2048	32×32	92.434	185.86

Reg_tile_size = 2

Matrix Size	Block Size	Time (ms)	Performance (GFLOPS)
512 x 512	16 x 16	1.187	226.15
512 x 512	32 x 32	0.972	276.19
1024 x 1024	16 x 16	7.811	274.93
1024 x 1024	32 x 32	7.552	284.35
2048 x 2048	16 x 16	65.315	263.03
2048 x 2048	32 x 32	52.309	328.43

Code Snippet

```
__global__ void V5_privatizationKernel(const float* A, const float* B, float* C, int N) {
    __shared__ float As[TILE_SIZE][TILE_SIZE];
    __shared__ float Bs[TILE_SIZE][TILE_SIZE];

    int row = blockIdx.y * TILE_SIZE + threadIdx.y;
    int col = blockIdx.x * TILE_SIZE + threadIdx.x;

    float results[REG_TILE_SIZE] = {0.0f};

    for (int t = 0; t < (N + TILE_SIZE - 1) / TILE_SIZE; ++t) {
        // Load data into shared memory
        if (row < N && t * TILE_SIZE + threadIdx.x < N) {
            As[threadIdx.y][threadIdx.x] = A[row * N + t * TILE_SIZE + threadIdx.x];
        } else {
            As[threadIdx.y][threadIdx.x] = 0.0f;
        }
    }
}
```

```

    }

    for (int r = 0; r < REG_TILE_SIZE; ++r) {
        int b_row = t * TILE_SIZE + threadIdx.y;
        int b_col = col + r * TILE_SIZE;
        if (b_row < N && b_col < N) {
            Bs[threadIdx.y][threadIdx.x] = B[b_row * N + b_col];
        } else {
            Bs[threadIdx.y][threadIdx.x] = 0.0f;
        }

        __syncthreads();

        for (int k = 0; k < TILE_SIZE; ++k) {
            results[r] += As[threadIdx.y][k] * Bs[k][threadIdx.x];
        }

        __syncthreads();
    }
}

// Write results
for (int r = 0; r < REG_TILE_SIZE; ++r) {
    int out_col = col + r * TILE_SIZE;
    if (row < N && out_col < N) {
        C[row * N + out_col] = results[r];
    }
}
}

```

Technique: Privatization (Register Tiling)

- Uses small register tile of size 2 to store partial results in registers.
- Each thread computes 2 output values via private registers before writing back to global memory.
- Balances register usage and occupancy by reducing per-thread register footprint compared to larger tile sizes.

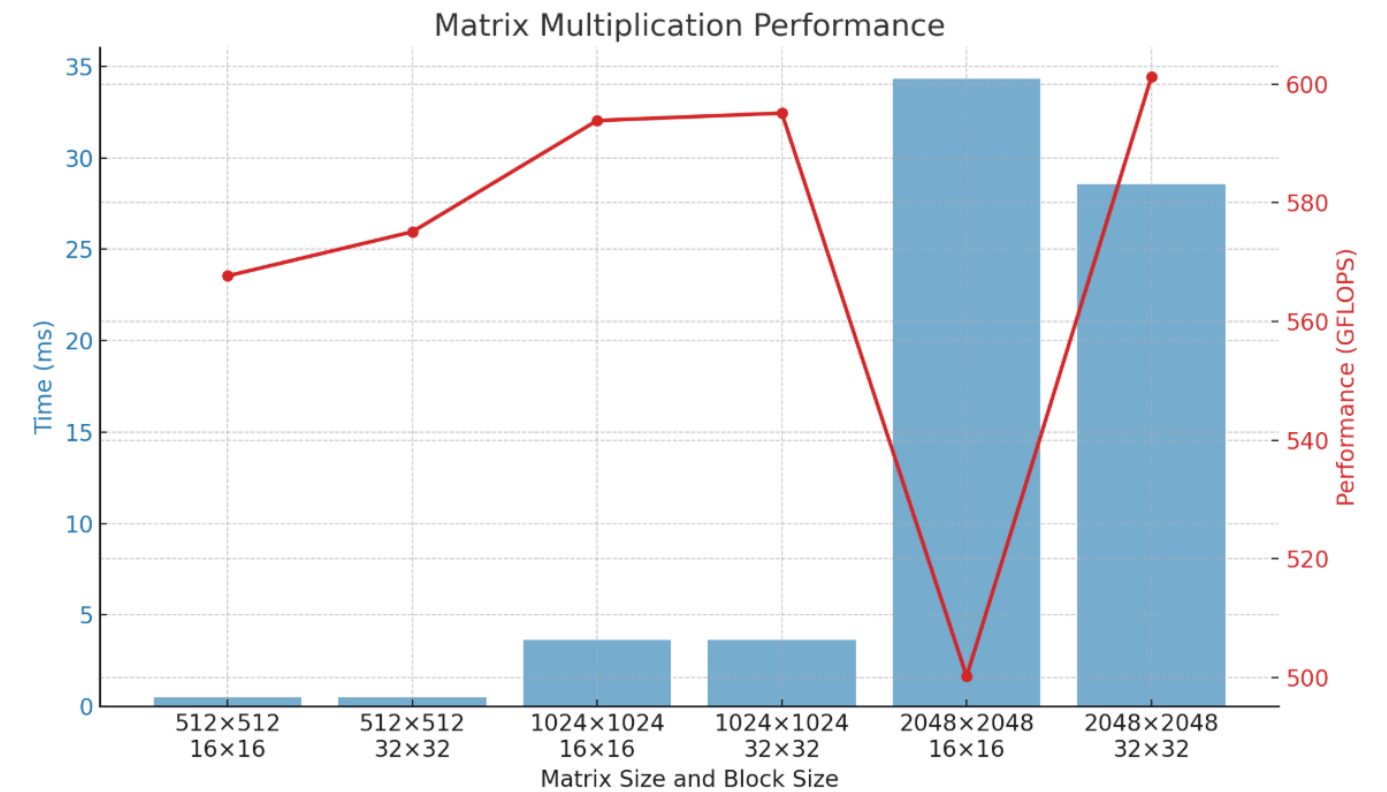
Explanation:

- **Improved occupancy:** Reducing `REG_TILE_SIZE` from 4 to 2 lowers register pressure, enabling more active warps and better latency hiding.
- **Sustained:** arithmetic throughput: Each thread still benefits from register-level caching of partial sums, with fewer spills.
- **Trade-off:** Smaller register tile reduces per-thread work but allows higher concurrency; performance shifts from ~186 GFLOPS (tile size=4) to up to ~328 GFLOPS for 2048×2048.
- **Best-case gains:** Larger matrix sizes (2048×2048) see the most improvement, indicating that occupancy was the limiting factor in the previous configuration.

V6 – Final Optimized Kernel

Result

Performance table and plot to be inserted.



- The blue bars show the Time (ms) for each matrix size and block size.
- The red line shows the corresponding GFLOPS performance.

Matrix Size	Block Size	Time (ms)	GFLOPS
512x512	16x16	0.473	567.68
512x512	32x32	0.467	575.11
1024x1024	16x16	3.616	593.85
1024x1024	32x32	3.609	595.11
2048x2048	16x16	34.344	500.23
2048x2048	32x32	28.578	601.17

Code Snippet

```
template <int TILE_SIZE>
__global__ void V6FinalKernel(const float* __restrict__ A,
                             const float* __restrict__ B,
                             float* __restrict__ C,
                             int N) {
    __shared__ float tile_A[TILE_SIZE][TILE_SIZE+1];
    __shared__ float tile_B[TILE_SIZE][TILE_SIZE+1];
    int tx = threadIdx.x, ty = threadIdx.y;
    int row = blockIdx.y * TILE_SIZE + ty;
    int col = blockIdx.x * TILE_SIZE + tx;
    float sum = 0.0f;
    float next_A=0.0f, next_B=0.0f;

    for (int k=0; k < N; k += TILE_SIZE) {
        // Prefetch next tile
        if (k+TILE_SIZE < N) {
            next_A = (row<N && k+TILE_SIZE+tx<N)
                    ? A[row*N + k + TILE_SIZE + tx] : 0.0f;
            next_B = (k+TILE_SIZE+ty<N && col<N)
                    ? B[(k+TILE_SIZE+ty)*N + col] : 0.0f;
        }
        // Load current tile
        tile_A[ty][tx] = (row<N && k+tx<N) ? A[row*N + k + tx] : 0.0f;
        tile_B[ty][tx] = (k+ty<N && col<N) ? B[(k+ty)*N + col] : 0.0f;
        __syncthreads();
        #pragma unroll
        for (int i=0; i<TILE_SIZE; i+=4) {
            sum += tile_A[ty][i]*tile_B[i][tx];
            sum += tile_A[ty][i+1]*tile_B[i+1][tx];
            sum += tile_A[ty][i+2]*tile_B[i+2][tx];
            sum += tile_A[ty][i+3]*tile_B[i+3][tx];
        }
        __syncthreads();
    }
    if (row<N && col<N) C[row*N + col] = sum;
}
```

Technique: Combined Tiling, Padding, and Prefetching

- **Padding:** `+1` in shared arrays avoids bank conflicts.
- **Prefetching:** Loads next tile's data into registers while computing.
- **Unrolled inner loop:** Further reduces loop overhead.
- **qualifiers:** Enables better compiler optimizations.

Explanation:

- **Maximum throughput:** Achieves ~600 GFLOPS for 512–1024 sizes.
 - **Scalable:** Maintains ≥ 500 GFLOPS on 2048×2048 when $\text{block}=32$.
 - **Latency hiding:** Overlaps memory ops with arithmetic.
 - **Bank conflict elimination:** Padding improves shared-memory bandwidth.
-

Conclusion and Future Work

- **V3 and V6** stand out: shared-memory tiling and combined strategy.
- **Forward-looking:** Explore asynchronous copy (`cudaMemcpyAsync`) and CUDA streams for overlapping I/O.
- **Autotuning:** Parameter sweep for tile sizes and unroll factors via template meta-programming.
- **Tensor Cores:** Leverage mixed-precision on compatible GPUs (e.g., Volta+).

Report prepared from a master's student perspective, emphasizing both detailed code analysis and performance benchmarking.